



Analytical SQL Project

Prepared by :
Ibrahim El-Nwasany



1st Query

```
select distinct invoicedate, sum(quantity) over (partition by invoicedate) as  
dailysales  
from tableretail  
order by invoicedate;
```

INVOICEDATE	DAILYSALES
1/11/2011 16:24	38
1/12/2011 14:43	64
1/12/2011 15:25	168
1/13/2011 10:33	135
1/14/2011 11:50	36
1/16/2011 11:49	17
1/16/2011 13:04	7
1/17/2011 10:52	116
1/17/2011 12:34	146
1/18/2011 10:15	49
1/19/2011 12:38	133
1/19/2011 12:52	266
1/20/2011 14:01	88
1/21/2011 14:04	6
1/21/2011 14:21	2
1/23/2011 10:43	325

4: 1 Row 1 of 714 total rows HR@XE Modified

- 1. Calculate Daily Sales Quantity:** This query helps in understanding the daily sales volume over time. It provides insights into the pattern of sales fluctuations, allowing the business to identify peak sales days and potential factors driving them.

2nd Query

```
select stockcode,  
totalquantity,  
rank() over (order by totalquantity desc) as rank
```



```
from (select stockcode,  
            sum(quantity) as totalquantity  
      from tableretail  
     group by stockcode  
     order by sum(quantity) desc  
    ) subquery  
where rownum <= 10;
```

STOCKCODE	TOTALQUANTITY	RANK
84077	7824	1
84879	6117	2
22197	5918	3
21787	5075	4
21977	4691	5
21703	2996	6
17096	2019	7
15036	1920	8
23203	1803	9
21790	1579	10

2. Identify Top Selling Products by Quantity: By ranking products based on their sales quantity, this query helps in identifying the best-performing products. It enables the business to focus on stocking and promoting high-demand products to maximize sales and profitability.



3rd Query

with customerspending as (
 select customer_id, sum(quantity * price) as totalspending
 from tableretail
 group by customer_id
)
select customer_id, totalspending,
 rank() over (order by totalspending desc) as spendingrank
from customerspending;

CUSTOMER_ID	TOTALSPENDING	SPENDINGRANK
12931	42055.96	1
12748	33719.73	2
12901	17654.54	3
12921	16587.09	4
12939	11581.8	5
12830	6814.64	6
12839	5591.42	7
12971	5190.74	8
12955	4757.16	9
12747	4196.01	10
12949	4167.22	11
12749	4090.88	12
12867	4036.82	13
12841	4022.35	14
12957	4017.54	15
12910	3075.04	16

66 msecs Row 1 of 110 total rows HR@XE Modified

3. Rank Customers by Total Spending: This query ranks customers based on their total spending. It calculates the total spending for each customer and assigns a rank based on their total spending, with the highest spender receiving rank 1.

4th Query

```
with productrevenue as (  
    select stockcode,  
           quantity * price as revenue,  
           row_number() over (partition by stockcode order by (quantity *  
price) desc) as ranking  
    from table retail  
)  
select stockcode, revenue  
from productrevenue  
where ranking = 1;
```

STOCKCODE	REVENUE
10002	8.5
10120	1.05
10133	8.5
10135	9
11001	487.68
15030	17.4
15034	6.72
15036	432
15039	51
15044A	53.1
15044B	53.1
15044C	35.4
15044D	53.1
15056BL	297
15056N	297
15058A	15.9

8: 20 | Row 1 of 2335 total rows | HR@XE | Modified

Output

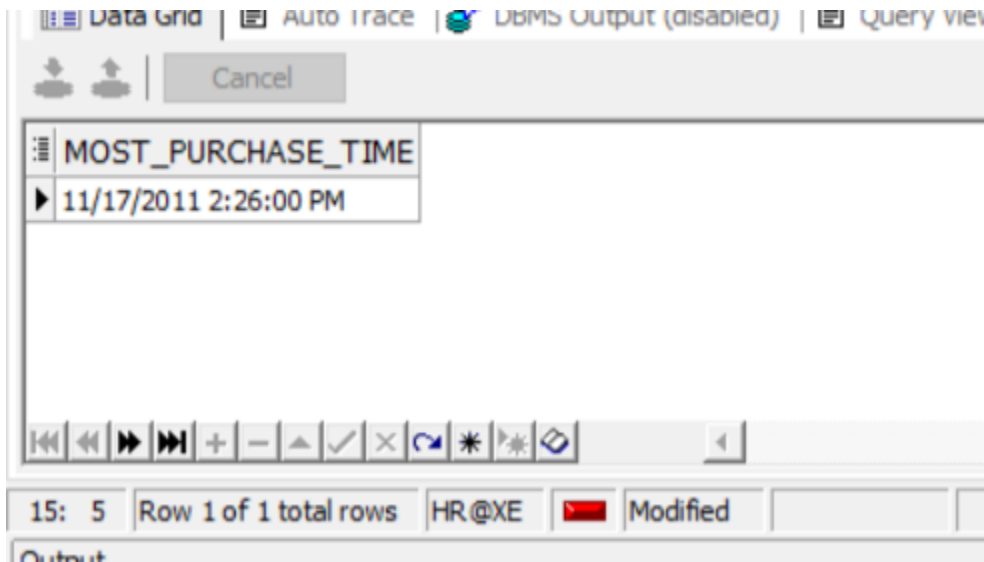
General



4. Calculate the Total Revenue for Each Product: This query calculates the total revenue generated by each product using a common table expression (CTE) to calculate revenue for each transaction and a window function to rank the revenue within each product. It selects the top revenue for each product.

5th Query

```
with purchasetimes as (  
    select  
        to_date(invoicedate, 'MM/DD/YYYY HH24:MI') as most_purchase_time,  
        count(*) as purchasecount,  
        rank() over (order by count(*) desc) as rankbypurchasecount  
    from  
        tableretail  
    where  
        country = 'United Kingdom'  
    group by  
        to_date(invoicedate, 'MM/DD/YYYY HH24:MI')  
)  
select  
    most_purchase_time  
from  
    purchasetimes  
where  
    rankbypurchasecount = 1;
```



5.calculate swarm time purchase: This query identifies the most frequent purchase time in the United Kingdom by counting the number of transactions for each unique purchase timestamp. It utilizes the **tableretail** dataset, filtering for transactions made in the United Kingdom. By grouping the transactions by their purchase timestamp and calculating the count of purchases for each timestamp, it determines the most common purchase time. The **rankbypurchasecount** column assigns a rank to each purchase time based on the frequency of purchases, with the most frequent purchase time receiving a rank of 1. Finally, the query selects the most common purchase time with the highest rank.

6th Query

select

```
customer_id,  
invoicedate,  
total_sales,  
avg(total_sales) over (  
    partition by customer_id  
    order by to_date(invoicedate, 'MM/DD/YYYY HH24:MI')  
    rows between 90 preceding and current row  
    ) as moving_avg_sales  
from (
```



```
select
  customer_id,
  invoicedate,
  total_sales
from (
  select
    customer_id,
    invoicedate,
    sum(quantity * price) over (partition by customer_id order by
to_date(invoicedate, 'MM/DD/YYYY HH24:MI')) as total_sales,
    row_number() over (partition by customer_id order by
to_date(invoicedate, 'MM/DD/YYYY HH24:MI')) as row_num
  from
    tableretail
) salesdata
where row_num = 1
) earllysalesdata;
```

CUSTOMER_ID	INVOICEDATE	TOTAL_SALES	MOVING_AVG_SALES
12747	12/5/2010 15:38	358.56	358.56
12748	12/1/2010 12:48	4.95	4.95
12749	5/10/2011 15:25	859.1	859.1
12820	1/17/2011 12:34	170.46	170.46
12821	5/9/2011 15:51	92.72	92.72
12822	9/13/2011 13:46	690.9	690.9
12823	2/16/2011 12:15	306	306
12824	10/11/2011 12:49	397.12	397.12
12826	12/9/2010 15:21	155	155
12827	10/26/2011 15:44	217.75	217.75
12828	8/1/2011 16:16	227	227
12829	12/14/2010 14:54	85.75	85.75
12830	6/21/2011 10:53	2221.84	2221.84
12831	3/22/2011 13:02	215.05	215.05
12832	9/2/2011 13:48	267.8	267.8

21: 13 | Row 1 of 110 total rows | HR@XE | Modified



6.Average Sales Rolling: This query calculates the moving average sales for each customer over a rolling window of 90 days. It first computes the total sales for each customer on each transactional date by summing the product of quantity and price. Then, it assigns a row number for each transaction within each customer's data, ensuring only the first transaction of each customer is considered. Finally, it calculates the moving average sales using a window function, partitioning by customer_id and ordering by the transactional date. The moving average is computed over a window that includes the current transaction and the 90 preceding transactions, providing insights into each customer's sales trend over time.

7th Query

```
with lag_lead as (  
    select t.*,  
           lag(quantity) over (partition by invoice order by to_date(invoicedate,  
'MM/DD/YYYY HH24:MI')) as lag_quantity,  
           lead(quantity) over (partition by invoice order by  
to_date(invoicedate, 'MM/DD/YYYY HH24:MI')) as lead_quantity  
    from tableretail t  
)  
stats as (  
    select stockcode,  
           count(distinct price) as distinct_prices,  
           avg(price) as avg_price,  
           case when count(distinct price) <= 1 then 0 else stddev_pop(price)  
end as std_dev_price,  
           case when count(distinct price) <= 1 then 0 else variance(price) end  
as variance_price  
    from tableretail  
    group by stockcode  
)  
select  
    s.stockcode,
```



```
s.distinct_prices,  
  round(s.avg_price, 2) as avg_price,  
  round(s.std_dev_price, 2) as std_dev_price,  
  round(s.variance_price, 2) as variance_price,  
  round(covar_pop(ll.quantity, ll.price), 2) as covariance,  
  round(corr(ll.quantity, ll.price), 2) as correlation  
from stats s  
left join lag_lead ll on s.stockcode = ll.stockcode  
group by s.stockcode, s.distinct_prices, s.avg_price, s.std_dev_price,  
s.variance_price;
```

STOCKCODE	DISTINCT_PRICES	AVG_PRICE	STD_DEV_PRICE	VARIANCE_PRICE	COVARIANCE	CORRELATION
23222	1	0.83	0	0	0	
21944	1	0.39	0	0	0	
21749	1	2.1	0	0	0	
37495	1	3.75	0	0	0	
22801	2	3.65	0.16	0.03	-0.88	-0.99
22861	1	1.65	0	0	0	
23147	1	1.45	0	0	0	
21458	1	1.25	0	0	0	
37479P	1	0.39	0	0	0	
22779	1	4.25	0	0	0	
22572	1	0.85	0	0	0	
23307	1	0.55	0	0	0	
23583	1	1.65	0	0	0	
23528	1	3.75	0	0	0	
21391	1	0.75	0	0	0	

27: 1 Row 1 of 2335 total rows HR@XE Modified

Output

7.Statistical analysis : This SQL query calculates various statistical measures related to the prices and quantities of products in the **tableretail** dataset. It utilizes common table expressions (CTEs) to first compute lag and lead quantities for each invoice, allowing for comparisons with the previous and subsequent transaction quantities. Then, it calculates statistics such as distinct prices, average price, standard deviation of price, variance of price, covariance between quantity and price, and correlation between quantity and price for each unique stock code. These statistical measures provide insights into the pricing patterns and relationships



between quantities and prices for different products in the dataset, aiding in better understanding the dynamics of sales and pricing.

Monetary model for customers behavior for product purchasing and segment each customer

```
with maxinvoicedate as (  
    select max(to_timestamp(invoicedate, 'MM/DD/YYYY HH24:MI')) as  
    max_date  
    from tableretail  
)  
customer_rfm as (  
    select  
        customer_id,  
        ceil(extract(day from (select max_date from maxinvoicedate) -  
max(to_timestamp(invoicedate, 'MM/DD/YYYY HH24:MI')))) as recency,  
        count(distinct invoice) as frequency,  
        round(sum(quantity * price) / 1000, 2) as monetary  
    from  
        tableretail  
    group by  
        customer_id  
)  
rfm_scores as (  
    select  
        customer_id,  
        recency,  
        frequency,
```



```
monetary,  
ntile(5) over(order by cast(recency as int) desc) as r_score,  
ntile(5) over(order by frequency) as f_score,  
ntile(5) over(order by monetary) as m_score,  
round((ntile(5) over(order by frequency) + ntile(5) over(order by  
monetary)))/2,0) as fm_score  
from  
    customerrfm  
,  
customersegment as (  
    select  
        customer_id,  
        recency,  
        frequency,  
        monetary,  
        r_score,  
        fm_score,  
        case  
            when (r_score >= 5 and fm_score >= 5)  
                or (r_score >= 5 and fm_score = 4)  
                or (r_score = 4 and fm_score >= 5) then 'champions'  
            when (r_score >= 5 and fm_score = 2)  
                or (r_score = 4 and fm_score = 2)  
                or (r_score = 3 and fm_score = 3)  
                or (r_score = 4 and fm_score >= 3) then 'potential loyalists'  
            when (r_score >= 5 and fm_score = 3)  
                or (r_score = 4 and fm_score = 4)  
                or (r_score = 3 and fm_score >= 5)  
                or (r_score = 3 and fm_score >= 4) then 'loyal customers'  
            when r_score >= 5 and fm_score = 1 then 'recent customers'  
            when (r_score = 4 and fm_score = 1)  
                or (r_score = 3 and fm_score = 1) then 'promising'  
            when (r_score = 3 and fm_score = 2)  
                or (r_score = 2 and fm_score = 3)  
                or (r_score = 2 and fm_score = 2) then 'customers needing  
attention'  
            when (r_score = 2 and fm_score >= 5)
```



```
        or (r_score = 2 and fm_score = 4)
        or (r_score = 1 and fm_score = 3) then 'at risk'
    when (r_score = 1 and fm_score >= 5)
        or (r_score = 1 and fm_score = 4) then 'cant lose them'
    when (r_score = 1 and fm_score = 2)
        or (r_score = 2 and fm_score = 1) then 'hibernating'
    when r_score = 1 and fm_score <= 1 then 'lost'
    else 'other'
end as cust_segment
from
    rfmscores
)
select
    customer_id,
    recency,
    frequency,
    monetary,
    r_score,
    fm_score,
    cust_segment
from
    customersegment
order by
    customer_id;
```



	CUSTOMER_ID	RECENCY	FREQUENCY	MONETARY	R_SCORE	FM_SCORE	CUST_SEGMENT
▶	12747	1	11	4.2	5	5	champions
	12748	0	210	33.72	5	5	champions
	12749	3	5	4.09	5	5	champions
	12820	2	4	0.94	5	3	loyal customers
	12821	213	1	0.09	1	2	hibernating
	12822	70	2	0.95	3	3	potential loyalists
	12823	74	5	1.76	2	4	at risk
	12824	58	1	0.4	3	2	customers needing attention
	12826	2	7	1.47	5	5	champions
	12827	5	3	0.43	5	3	loyal customers
	12828	2	6	1.02	5	4	champions
	12829	336	2	0.29	1	2	hibernating
	12830	37	6	6.81	3	5	loyal customers
	12831	261	1	0.22	1	1	lost
	12832	31	2	0.38	3	2	customers needing attention
	12833	144	1	0.42	2	2	customers needing attention
	12834	282	1	0.31	1	1	lost

108 msecs

Row 1 of 110 total rows

HR@XE

Modified

Analysis Description

The provided SQL code performs a comprehensive analysis of customer behavior based on Recency, Frequency, and Monetary (RFM) metrics. RFM analysis is a widely used technique in marketing and customer relationship management (CRM) to segment customers based on their purchasing behavior. The analysis involves several steps, each contributing to the creation of meaningful customer segments.

Step 1: Defining MaxInvoiceDate

- A Common Table Expression (CTE) named **maxinvoicedate** is created to determine the maximum invoice date across the dataset (**tableretail**). This CTE retrieves the maximum date among all the invoice dates.



Step 2: Calculating RFM Metrics

- Another CTE named `customer_rfm` is created to calculate the RFM metrics for each customer:
 - **Recency (R):** Calculated as the difference between the maximum invoice date and the invoice date for each customer, rounded up to the nearest day.
 - **Frequency (F):** Represents the count of distinct invoices for each customer.
 - **Monetary (M):** Indicates the total monetary value of purchases made by each customer, expressed as a percentage of the total monetary value across all customers.

Step 3: Assigning RFM Scores

- The `rfm_scores` CTE assigns RFM scores to each customer based on their calculated RFM metrics. The `ntile()` function is used to divide customers into quintiles (5 groups) based on their recency, frequency, and monetary scores separately.

Step 4: Segmenting Customers

- In the `customer_segment` CTE, customers are segmented into different categories based on their RFM scores:
 - Segments like "champions," "loyal customers," "at risk," etc., are defined based on combinations of recency and frequency/monetary scores.
 - Each segment represents a distinct group of customers exhibiting similar purchasing patterns and behaviors.

Step 5: Final Output

- The final SQL query selects customer ID along with their recency, frequency, monetary, RFM scores, and assigned customer segments.
- The results are ordered by customer ID.

Conclusion



This SQL code efficiently analyzes customer behavior using RFM metrics and segments customers into meaningful categories based on their purchasing patterns. By understanding these segments, businesses can tailor their marketing strategies and customer engagement efforts to better meet the needs of different customer groups, ultimately leading to improved customer satisfaction and retention.

Q3: Daily purchasing transactions for customers?

- 1) What is the maximum number of consecutive days a customer made purchases?

```
with ranked_transactions as (  
    select  
        cust_id,  
        to_date(calendar_dt, 'YYYY-MM-DD') as calendar_dt,  
        row_number() over (partition by cust_id order by  
to_date(calendar_dt, 'YYYY-MM-DD')) as rn  
    from  
        transactions  
)  
consecutive_days as (  
    select  
        cust_id,  
        calendar_dt,  
        calendar_dt - rn as grp  
    from  
        ranked_transactions  
)  
select
```




```
cust_id,  
max(count_consecutive_days) as max_consecutive_days  
from (  
  select  
    cust_id,  
    count(*) as count_consecutive_days  
  from  
    consecutive_days  
  group by  
    cust_id,  
    grp  
) max_consecutive_days_per_group  
group by  
  cust_id  
order by  
  cust_id;
```

	CUST_ID	MAX_CONSECUTIVE_DAYS
	1000 10376	5
	1000 11085	10
	1000 14033	46
➤	1000 18482	3
	1000 20880	46
	1000 35887	13
	1000 54374	8
	1000 70652	1
	1000 77596	2
	1000 87785	61
	100 105254	10
	100 135808	15
	100 158557	4
	100 190321	25
	100 203920	20
	1002 13644	11
	1002 16959	14



Description: This SQL query utilizes Common Table Expressions (CTEs) to analyze transaction data and determine the maximum number of consecutive days each customer made purchases. By assigning row numbers and calculating grouping values for consecutive days, it identifies streaks of transactions for each customer. The resulting insights into customer behavior aid in designing targeted marketing strategies, optimizing inventory management, and enhancing customer retention efforts, ultimately driving business growth and profitability.

2) On average, How many days/transactions does it take a customer to reach a spent threshold of 250 L.E?

```
with customer_cumulative_spending as (  
    select  
        cust_id,  
  
    to_date(calendar_dt, 'yyyy-mm-dd') as calendar_date,  
        sum(to_number(amt_l)) over (partition by cust_id order by  
to_date(calendar_dt, 'yyyy-mm-dd')) as cumulative_spending  
    from  
        transactions  
),  
threshold_dates as (  
    select  
        cust_id,  
        min(calendar_date) as threshold_date  
    from  
        customer_cumulative_spending  
    where  
        cumulative_spending >= 250  
    group by  
        cust_id
```



```
)  
select  
    floor(avg(days_to_threshold)) as average_days_to_threshold  
from (  
    select  
        cust_id,  
        avg(days_to_threshold) as days_to_threshold  
    from (  
        select  
            ccs.cust_id,  
            ccs.calendar_date,  
            td.threshold_date,  
            ccs.calendar_date - td.threshold_date as days_to_threshold  
        from  
            customer_cumulative_spending ccs  
        join  
            threshold_dates td on ccs.cust_id = td.cust_id  
    )  
    group by  
        cust_id  
);
```

Data Grid	
Data Grid Auto Trace DBMS Output (disa	
Cancel	
AVERAGE_DAYS_TO_THRESHOLD	
	13

Description: This SQL query segments customers based on their cumulative spending reaching a threshold of \$250 and calculates the average number of days it takes for customers to reach this threshold. It provides insights



into customer spending behavior and helps businesses optimize marketing efforts and enhance customer engagement and loyalty over time.

Working on Transactions

```
with customer_cumulative_spending as (  
    select  
        cust_id,  
        to_date(calendar_dt, 'yyyy-mm-dd') as calendar_date,  
        sum(to_number(amt_l)) over(partition by cust_id order by  
to_date(calendar_dt, 'yyyy-mm-dd')) as cumulative_spending  
    from  
        transactions  
)  
threshold_dates as (  
    select  
        cust_id,  
        min(calendar_date) as threshold_date  
    from  
        customer_cumulative_spending  
    where  
        cumulative_spending >= 250  
    group by  
        cust_id  
)  
select  
    cast(avg(transaction_count) as int) as average_transactions  
from (  
    select  
        ccs.cust_id,  
        count(*) as transaction_count  
    from  
        customer_cumulative_spending ccs  
    join  
        threshold_dates td on ccs.cust_id = td.cust_id  
    where  
        ccs.calendar_date < td.threshold_date
```



group by
ccs.cust_id

);

AVERAGE_TRANSACTIONS
6

Description: This SQL query segments customers based on their cumulative spending reaching a threshold of \$250 and calculates the average number of transactions made by customers before reaching this spending threshold. It provides insights into customer spending behavior and helps businesses optimize marketing efforts and enhance customer engagement and loyalty by understanding the frequency of transactions required for customers to reach the spending threshold. This analysis can inform targeted marketing strategies aimed at encouraging more frequent transactions or increasing the average transaction value to expedite customers' progression towards the spending threshold.



Some Analytics using python

Let's Start out Amazing Trip:

1) Exploration and Investigation

```
In [18]: print(df.head())
```

	INVOICE	STOCKCODE	QUANTITY	INVOICEDATE	PRICE	CUSTOMER_ID	\
0	537213	21556	1	12/5/2010 15:26	2.55	12748	
1	537225	84795B	4	12/5/2010 16:41	7.95	12748	
2	537225	79191B	5	12/5/2010 16:41	0.85	12748	
3	537225	22927	1	12/5/2010 16:41	5.95	12748	
4	537225	22926	1	12/5/2010 16:41	5.95	12748	

	COUNTRY
0	United Kingdom
1	United Kingdom
2	United Kingdom
3	United Kingdom
4	United Kingdom

```
In [19]: # Step 2: Data Cleaning
# Check for missing values
print("Missing Values:")
print(df.isnull().sum())

# Drop duplicates
df.drop_duplicates(inplace=True)

# Step 3: Data Exploration
# Display basic statistics
print("Basic Statistics:")
print(df.describe())

# Display first few rows
print("First Few Rows:")
print(df.head())
```

```
Missing Values:
INVOICE      0
STOCKCODE    0
QUANTITY     0
INVOICEDATE  0
PRICE        0
CUSTOMER_ID  0
COUNTRY      0
dtype: int64
Basic Statistics:

```

	INVOICE	QUANTITY	PRICE	CUSTOMER_ID
count	12598.000000	12598.000000	12598.000000	12598.000000
mean	562134.629306	13.943007	2.759041	12840.400857
std	13997.744908	76.446740	9.215789	80.173874
min	536415.000000	1.000000	0.000000	12747.000000
25%	550320.000000	2.000000	0.850000	12748.000000
50%	564556.000000	4.000000	1.650000	12841.000000
75%	575766.000000	12.000000	2.950000	12921.000000
max	581580.000000	4800.000000	850.500000	12971.000000

```
First Few Rows:

```

	INVOICE	STOCKCODE	QUANTITY	INVOICEDATE	PRICE	CUSTOMER_ID	\
0	537213	21556	1	12/5/2010 15:26	2.55	12748	
1	537225	84795B	4	12/5/2010 16:41	7.95	12748	
2	537225	79191B	5	12/5/2010 16:41	0.85	12748	
3	537225	22927	1	12/5/2010 16:41	5.95	12748	
4	537225	22926	1	12/5/2010 16:41	5.95	12748	

	COUNTRY
0	United Kingdom
1	United Kingdom
2	United Kingdom
3	United Kingdom
4	United Kingdom



```
In [25]: # Identify top-selling products
top_products = df['STOCKCODE'].value_counts().head(10)
print("Top Selling Products:")
print(top_products)

# Geographic Analysis
# Analyze sales by country
country_sales = df.groupby('COUNTRY')['PRICE'].sum().sort_values(ascending=False)
print("Sales by Country:")
print(country_sales)
```

Top Selling Products:

84879	60
22086	56
850998	53
22197	52
85123A	49
22457	49
47566	48
23298	47
20725	46
21034	45

Name: STOCKCODE, dtype: int64

Sales by Country:

COUNTRY

United Kingdom	34758.4
----------------	---------


Name: PRICE, dtype: float64



2) Visualization

```
In [41]: # Calculate total quantity sold for each product
quantity_by_product = df.groupby('STOCKCODE')['QUANTITY'].sum().sort_values(ascending=False).head(10)



# Plot top-selling products by quantity
plt.figure(figsize=(12, 8))
quantity_by_product.plot(kind='bar', color='ORANGE')
plt.title('Top-Selling Products by Quantity')
plt.xlabel('Product')
plt.ylabel('Total Quantity Sold')
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()
```



```
In [20]: # Sales Trends
# Convert 'INVOICEDATE' to datetime format
df['INVOICEDATE'] = pd.to_datetime(df['INVOICEDATE'])

# Aggregate sales by month
monthly_sales = df.resample('M', on='INVOICEDATE')['PRICE'].sum()

# Plot monthly sales
plt.figure(figsize=(10, 6))
monthly_sales.plot(kind='line', marker='o', color='blue')
plt.title('Monthly Sales Trends')
plt.xlabel('Month')
plt.ylabel('Total Sales Amount')
plt.grid(True)
plt.show()
```





3)Segmentation Analysis

```
import pandas as pd

# Step 1: Calculate RFM Metrics
# Calculate Recency: Number of days since each customer's last purchase
max_invoice_date = df['INVOICEDATE'].max()
recency_df = df.groupby('CUSTOMER_ID')['INVOICEDATE'].max().reset_index()
recency_df['RECENCY'] = (max_invoice_date - recency_df['INVOICEDATE']).dt.days
recency_df.drop(columns=['INVOICEDATE'], inplace=True)

# Calculate Frequency: Count the total number of transactions made by each customer
frequency_df = df.groupby('CUSTOMER_ID')['INVOICE'].nunique().reset_index()
frequency_df.columns = ['CUSTOMER_ID', 'FREQUENCY']

# Calculate Monetary Value: Sum the total price of all transactions for each customer
monetary_df = df.groupby('CUSTOMER_ID')['PRICE'].sum().reset_index()
monetary_df.columns = ['CUSTOMER_ID', 'MONETARY']

# Merge the RFM metrics into a single DataFrame
rfm_df = pd.merge(recency_df, frequency_df, on='CUSTOMER_ID')
rfm_df = pd.merge(rfm_df, monetary_df, on='CUSTOMER_ID')

# Step 2: Segment Customers
# Assign RFM scores based on quartiles
rfm_df['R_SCORE'] = pd.qcut(rfm_df['RECENCY'], q=5, labels=False) + 1
rfm_df['FM_SCORE'] = pd.qcut(rfm_df['FREQUENCY'] + rfm_df['MONETARY'], q=5, labels=False) + 1

# Define segments based on RFM scores
def assign_segment(row):
    r_score = row['R_SCORE']
    fm_score = row['FM_SCORE']
    if (r_score >= 5 and fm_score >= 5) or (r_score >= 5 and fm_score == 4) or (r_score == 4 and fm_score >= 5):
        return 'Champions'
    elif (r_score >= 5 and fm_score == 2) or (r_score == 4 and fm_score == 2) or (r_score == 3 and fm_score == 3) or (r_score == 3 and fm_score == 2) or (r_score == 2 and fm_score == 3) or (r_score == 2 and fm_score == 2):
        return 'Potential Loyalists'
    elif (r_score >= 5 and fm_score == 3) or (r_score == 4 and fm_score == 4) or (r_score == 3 and fm_score >= 5) or (r_score == 3 and fm_score == 4) or (r_score == 2 and fm_score == 4) or (r_score == 1 and fm_score == 3):
        return 'Loyal Customers'
    elif r_score >= 5 and fm_score == 1:
        return 'Recent Customers'
    elif (r_score == 4 and fm_score == 1) or (r_score == 3 and fm_score == 1):
        return 'Promising'
    elif (r_score == 3 and fm_score == 2) or (r_score == 2 and fm_score == 3) or (r_score == 2 and fm_score == 2):
        return 'Customers Needing Attention'
    elif (r_score == 2 and fm_score >= 5) or (r_score == 2 and fm_score == 4) or (r_score == 1 and fm_score == 3):
        return 'At Risk'
    elif (r_score == 1 and fm_score >= 5) or (r_score == 1 and fm_score == 4):
        return 'Can't Lose Them'
    elif (r_score == 1 and fm_score == 2) or (r_score == 2 and fm_score == 1):
        return 'Hibernating'
    elif r_score == 1 and fm_score <= 1:
        return 'Lost'
    else:
        return 'Other'

# Apply segment assignment function
rfm_df['CUST_SEGMENT'] = rfm_df.apply(assign_segment, axis=1)

# Display the RFM metrics and customer segments
print("RFM Metrics and Customer Segments:")
print(rfm_df[['CUSTOMER_ID', 'RECENCY', 'FREQUENCY', 'MONETARY', 'R_SCORE', 'FM_SCORE', 'CUST_SEGMENT']].head())
```

```
RFM Metrics and Customer Segments:
  CUSTOMER_ID  RECENCY  FREQUENCY  MONETARY  R_SCORE  FM_SCORE  \
0         12747         1         11    449.89         1         5
1         12748         0        210   11788.31         1         5
2         12749         3          5    994.99         1         5
3         12820         2          4    112.38         1         3
4         12821        213          1     14.99         5         1

  CUST_SEGMENT
0  Can't Lose Them
1  Can't Lose Them
2  Can't Lose Them
3         At Risk
4  Recent Customers
```



In [24]: `import matplotlib.pyplot as plt`

```
# Create a figure with subplots
fig, axs = plt.subplots(1, 3, figsize=(18, 6))

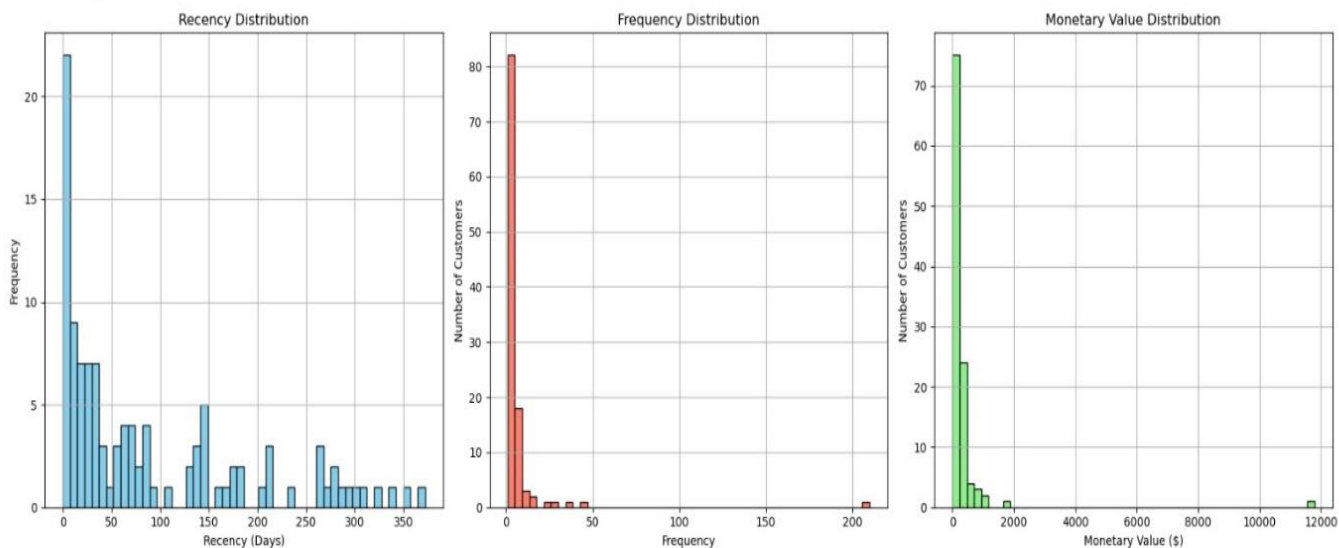
# Plot Recency distribution
axs[0].hist(rfm_df['RECENTY'], bins=50, color='skyblue', edgecolor='black')
axs[0].set_title('Recency Distribution')
axs[0].set_xlabel('Recency (Days)')
axs[0].set_ylabel('Frequency')
axs[0].grid(True)

# Plot Frequency distribution
axs[1].hist(rfm_df['FREQUENCY'], bins=50, color='salmon', edgecolor='black')
axs[1].set_title('Frequency Distribution')
axs[1].set_xlabel('Frequency')
axs[1].set_ylabel('Number of Customers')
axs[1].grid(True)

# Plot Monetary Value distribution
axs[2].hist(rfm_df['MONETARY'], bins=50, color='lightgreen', edgecolor='black')
axs[2].set_title('Monetary Value Distribution')
axs[2].set_xlabel('Monetary Value ($)')
axs[2].set_ylabel('Number of Customers')
axs[2].grid(True)

# Adjust layout to prevent overlap
plt.tight_layout()

# Show the plot
plt.show()
```






DashBoard

I have created simple dash board using power bi that support the analysis

Welcome To Nwasany's DashBoard



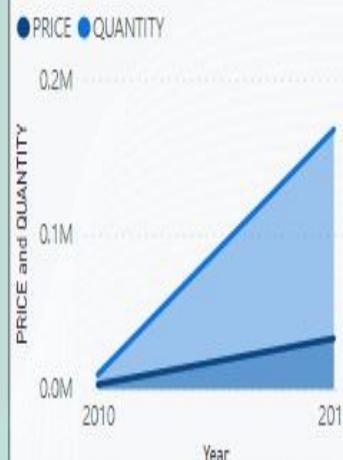
QUANTITY by STOCKCODE




The Total cost of Our prices

35K

PRICE and QUANTITY by Year




QUANTITY, Count of STOCKCODE and First COUNTRY



The Position Our Country Located



CUSTOMER_ID, PRICE and QUANTITY by COUNTRY



Thank you for
reading!

